Eunice Olorunshola

10/28/2021

CSC 6850 Machine Learning

Report

Note: I used the link provided for implementing linear regression and also from k fold link on the same website to help with this homework

1) The way that I did this assisgnment is that first step I created my dataset using the 3 links provided and turned it into a csv file than from there I implemented it into my code so it could read the csv file and

Screenshot of the output :

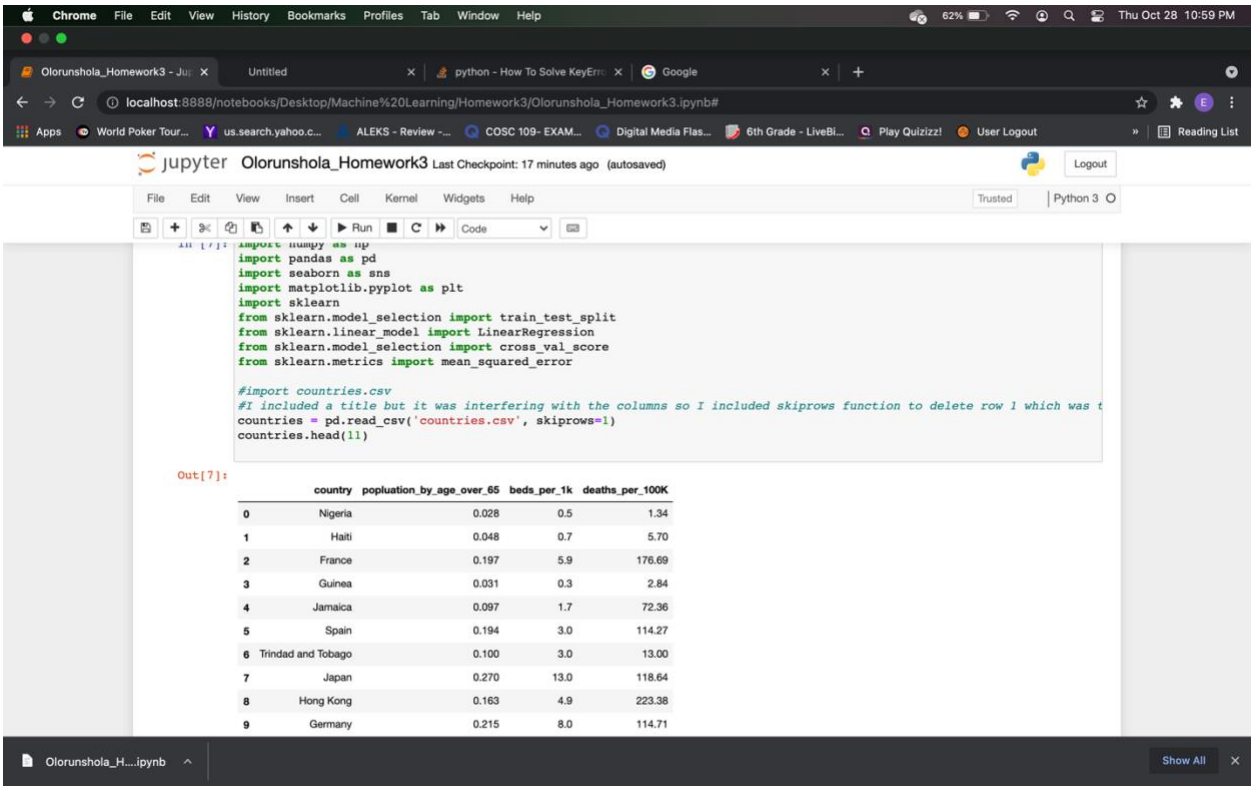

2) Second step is that I used the link provided to implement the linear regression I also used the k fold link from the same website scikit-learn.org to help implement linear regression below is the output of the results from implementing linear regression

Screenshot of the output:

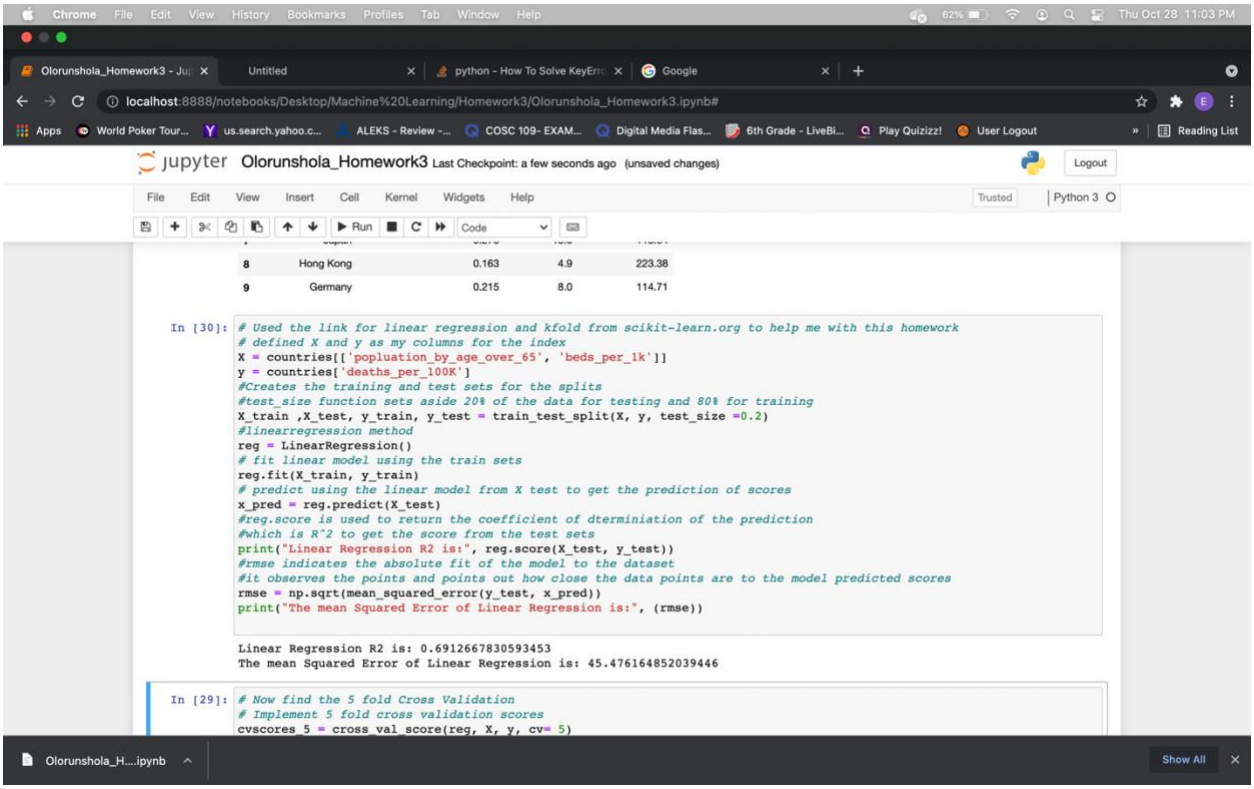

8	Hong Kong	0.163	4.9	223.38
9	Germany	0.215	8.0	114.71

In [30]:
```
# Used the link for linear regression and kfold from scikit-learn.org to help me with this homework
# defined X and y as my columns for the index
X = countries[['popluation_by_age_over_65', 'beds_per_1k']]
y = countries['deaths_per_100K']
#Creates the training and test sets for the splits
#test_size function sets aside 20% of the data for testing and 80% for training
X_train ,X_test, y_train, y_test = train_test_split(X, y, test_size =0.2)
#linearregression method
reg = LinearRegression()
# fit linear model using the train sets
reg.fit(X_train, y_train)
# predict using the linear model from X test to get the prediction of scores
x_pred = reg.predict(X_test)
#reg.score is used to return the coefficient of dterminiation of the prediction
#which is R^2 to get the score from the test sets
print("Linear Regression R2 is:", reg.score(X_test, y_test))
#rmse indicates the absolute fit of the model to the dataset
#it observes the points and points out how close the data points are to the model predicted scores
rmse = np.sqrt(mean_squared_error(y_test, x_pred))
print("The mean Squared Error of Linear Regression is:", (rmse))
```

```
Linear Regression R2 is: 0.6912667830593453
The mean Squared Error of Linear Regression is: 45.476164852039446
```

In [29]:
```
# Now find the 5 fold Cross Validation
# Implement 5 fold cross validation scores
cvscores_5 = cross_val_score(reg, X, y, cv= 5)
```

3) The last step is to implement 5 fold cross validation score from using the module cross_val_score from sk.learn I was able to implement the 5 scores from my dataset model correctly

Screenshot of output:

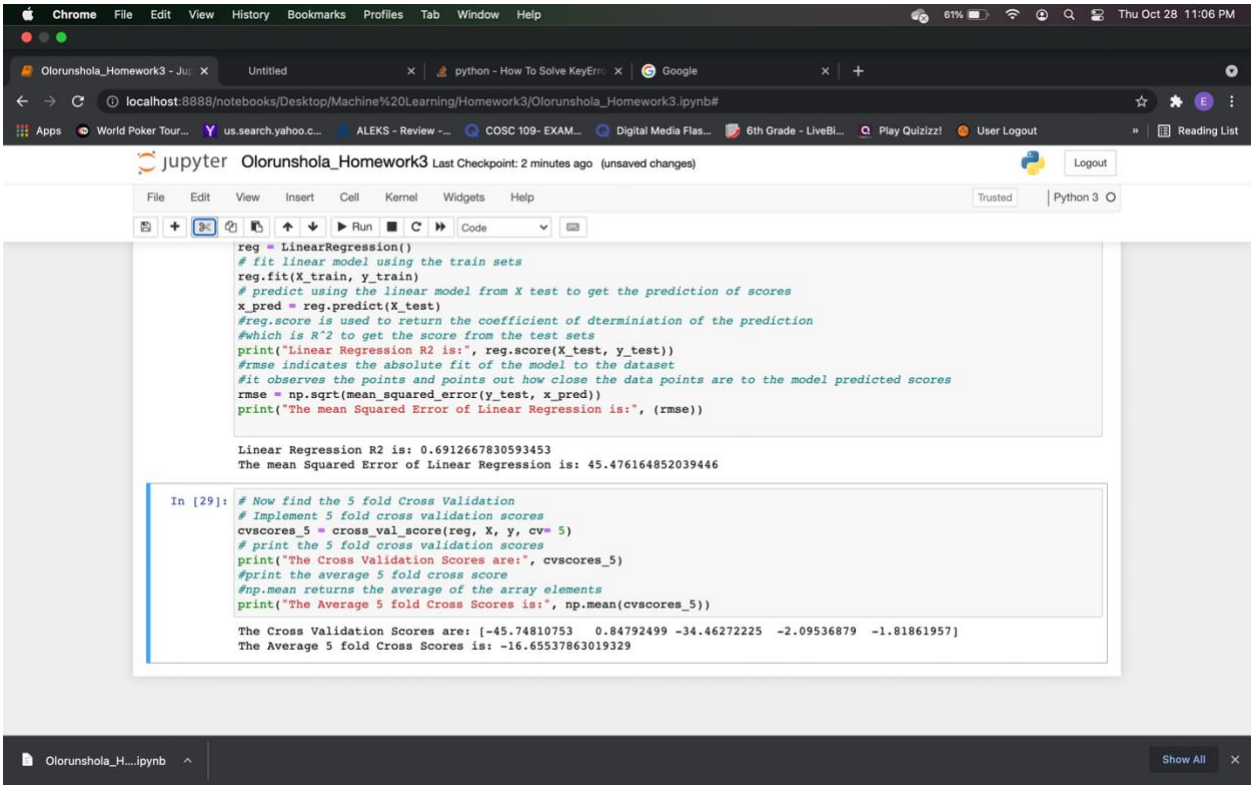

Appendix section

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error

#import countries.csv
#I included a title but it was interfering with the columns so I included skiprows function to delete row 1 which was the title
countries = pd.read_csv('countries.csv', skiprows=1)
countries.head(11)


# Used the link for linear regression and kfold from scikit-learn.org to help me with this homework
# defined X and y as my columns for the index
X = countries[['popluation_by_age_over_65', 'beds_per_1k']]
y = countries['deaths_per_100K']
#Creates the training and test sets for the splits
#test_size function sets aside 20% of the data for testing and 80% for training
X_train ,X_test, y_train, y_test = train_test_split(X, y, test_size =0.2)
#linearregression method
reg = LinearRegression()
# fit linear model using the train sets
reg.fit(X_train, y_train)
# predict using the linear model from X test to get the prediction of scores
x_pred = reg.predict(X_test)
#reg.score is used to return the coefficient of dterminiation of the prediction
#which is R^2 to get the score from the test sets
print("Linear Regression R2 is:", reg.score(X_test, y_test))
#rmse indicates the absolute fit of the model to the dataset
#it observes the points and points out how close the data points are to the model predicted scores
rmse = np.sqrt(mean_squared_error(y_test, x_pred))
print("The mean Squared Error of Linear Regression is:", (rmse))
```

```
#Now find the 5 fold Cross Validation
# Implement 5 fold cross validation scores
cvscores_5 = cross_val_score(reg, X, y, cv= 5)
# print the 5 fold cross validation scores
print("The Cross Validation Scores are:", cvscores_5)
#print the average 5 fold cross score
#np.mean returns the average of the array elements
print("The Average 5 fold Cross Scores is:", np.mean(cvscores_5))
```